

UNPACKING AN ANOMALY

Tales from the Ops Side



VM FARMS

@vmfarms

By Hany Fahim
Founder and CEO
@iHandroid



IT WAS A QUIET
NOVEMBER DAY

`“Content is not a valid RSA or DSA SSH Key.”`

AN EXCEPTION WAS THROWN

This error would haunt us for a month.

SSH KEY MANAGEMENT

- App in our portal to manage SSH keys.
- Allows one to map users and their SSH keys to servers.
- It also validates that the SSH key is valid and complete.

parse_sshkey()

```
1 try:  
2     data_b64 = base64.decodestring(key)  
3     int_len = 4  
4     str_len = struct.unpack('>I', data_b64[:int_len])[0]  
5     if str_len == 7:  
6         key_type = data_b64[int_len:int_len + str_len]  
7 except Exception:  
8     raise ValidationError('Contents is not a valid RSA or DSA SSH Key.')
```

WHY DID IT BREAK?

MEMORY CORRUPTION?

- We thought perhaps it was memory corruption, as it did not make sense to break here.
- Tried to replicate, but failed.
- We thought this was an anomaly and would **never happen again.**

“Content is not a valid RSA or DSA SSH Key.”

6 DAYS LATER

It happened again.

THIS DIDN'T MAKE MUCH SENSE

- Due to the generic exception handler, the real exception was being swallowed up.
- Very little to go on here.
- Decided to add some extra logging and wait for it to happen again.


```
1 import logging
2 logger = logging.getLogger(__name__)
3 ...
4 try:
5     data_b64 = base64.decodestring(key)
6     int_len = 4
7     str_len = struct.unpack('>I', data_b64[:int_len])[0]
8     if str_len == 7:
9         key_type = data_b64[int_len:int_len + str_len]
10 except Exception:
11     message = 'Content is not a valid RSA or DSA SSH Key.'
12     logger.exception(message)
13     raise ValidationError(message)
```

MOAR LOGGING

“Content is not a valid RSA or DSA SSH Key.”

6 DAYS LATER

It was back.

NOW WE KNOW THE ERROR

TypeError: `Struct()` argument `1` must be string, **not** `unicode`

FROM THIS LINE

```
7 str_len = struct.unpack('>I', data_b64[:int_len])[0]
```

But it still didn't make sense

MORE CONFUSION

- What was more confusing is it should never error at this stage.
- This function is used for both validation and parsing of the SSH Key.
- This error was being thrown **after** validation.
- We could not replicate the issue. We tried 1,000,000 times.

DEEPER DIVE

- Looking through the error log, it seemed to error out nearly 100 times.
- Noticed a strange pattern in the logs.
- Approximately 1 in 8 calls were failing.
- We have 8 workers spawned.
- Is it possible a Unicorn worker got corrupted somehow?

```
1 import logging
2 import os
3 logger = logging.getLogger(__name__)
4 ...
5 try:
6     data_b64 = base64.decodestring(key)
7     int_len = 4
8     str_len = struct.unpack('>I', data_b64[:int_len])[0]
9     if str_len == 7:
10         key_type = data_b64[int_len:int_len + str_len]
11 except Exception:
12     message = 'Content is not a valid RSA or DSA SSH Key.'
13     logger.exception(message, extra={'pid': os.getpid()})
14     raise ValidationError(message)
```

MOAR LOGGING

“Content is not a valid RSA or DSA SSH Key.”

7 DAYS LATER

It was back.

ASSUMPTION CONFIRMED!

- It was the same PID every time.
- A worker was getting corrupted. *Somehow.*
- Now what?

PROBLEM STOPPED.

SELF-HEALING POWERS

A faint, semi-transparent silhouette of a person sitting in a meditative lotus position is centered in the background. The background itself is a soft-focus image of a sunset or sunrise over a body of water, with a bright sun low on the horizon and scattered clouds.

- The problem appeared to resolve itself while troubleshooting.
- This is something we never noticed previously.
- Perhaps there is some previous request that causes the corruption?
- Needs moar logging!


```
1 def post_request(worker, req, environ, resp):
2     timestamp = datetime.now()
3
4     worker.log.info('{method} {uri} status:{status} req_no:{req_no}
(CORRUPTION LOGGING)'.format(
5         method=req.method,
6         req_no=worker.nr,
7         uri=req.uri,
8         status=resp.status_code,
9     ))
```

MOAR LOGGING!

“Content is not a valid RSA or DSA SSH Key.”

5 DAYS LATER

It was back.

SAME STORY

- Same PID.
- Slew of errors during a run, at an impossible stage.
- Stops on its own.
- Looked at our new logging.

GUNICORN LOGGING

- It did not appear to be any particular sequence of requests that caused the worker to break.
- When it did break, **all** calls to `parse_sshkey()` would break, for a short period of time, then fix itself **without** a restart.
- It did not appear to be related to the number of requests.



NOW WHAT?

WE WERE REACHING THE END

- We were out of ideas for debugging this further.
- Could not replicate this issue.
- Perhaps we should just replace calls to `struct` and abandon ship?


```
1 SSH_KEY_PREFIX = ['\x00', '\x00', '\x00', '\x07']
2
3 int_len = 4
4
5 try:
6     assert list(data_b64[0:int_len]) == SSH_KEY_PREFIX
7 except AssertionError:
8     message = 'Content is not a valid RSA or DSA SSH Key.'
9     logger.exception(message, extra={'pid': os.getpid()})
10    raise ValidationError(message)
```

NEW PR. SIMPLIFY.

No more `struct.unpack()`.
However, it was not merged.

“Content is not a valid RSA or DSA SSH Key.”

5 DAYS LATER

It was back.

BUG NEEDS TO BE SQUISHED

- PR hadn't been merged yet, but everyone was keen on resolving.
- Fresh eyes were on the problem.
- Important questions being asked.
- “Why does this ever work?”

```
1 from __future__ import unicode_literals
2 import logging
3 import os
4 logger = logging.getLogger(__name__)
5 ...
6 try:
7     data_b64 = base64.decodestring(key)
8     int_len = 4
9     str_len = struct.unpack('>I', data_b64[:int_len])[0]
10    if str_len == 7:
11        key_type = data_b64[int_len:int_len + str_len]
12 except Exception:
13     message = 'Content is not a valid RSA or DSA SSH Key.'
14     logger.exception(message, extra={'pid': os.getpid()})
15     raise ValidationError(message)
```

IT'S UNICODE!

EVERYTHING WAS UNICODE

- That import was added a few months back.
- It was part of a Django upgrade and an effort to be more Py3 compatible.
- This means...

THIS MEANS:

```
9 str_len = struct.unpack('>I', data_b64[:int_len])[0]
```

IS ACTUALLY:

```
9 str_len = struct.unpack(u'>I', data_b64[:int_len])[0]
```


WITH IMPORT

```
>>> from __future__ import unicode_literals
>>> import struct
>>> struct.unpack('>I', '\x00\x00\x00\x07')
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
TypeError: Struct() argument 1 must be string, not unicode
```

AND WITHOUT

```
>>> import struct
>>> struct.unpack('>I', '\x00\x00\x00\x07')
(7,)
```



BUT HOW DID THIS
EVER WORK?


```
$ ./manage.py shell_plus
```

```
...
```

```
Python 2.7.6 (default, Nov 26 2013, 13:28:29)
```

```
Type "copyright", "credits" or "license" for more information.
```

```
IPython 5.0.0 -- An enhanced Interactive Python.
```

```
? -> Introduction and overview of IPython's features.
```

```
%quickref -> Quick reference.
```

```
help -> Python's own help system.
```

```
object? -> Details about 'object', use 'object??' for extra details.
```

```
In [1]: from __future__ import unicode_literals
```

```
In [2]: import struct
```

```
In [3]: struct.unpack('>I', '\x00\x00\x00\x07')
```

```
Out[3]: (7,)
```

```
In [4]:
```

TRIED TO RUN WITHIN DJANGO

```
$ ./manage.py shell --plain
Python 2.7.6 (default, Nov 26 2013, 13:28:29)
[GCC 4.4.6 20120305 (Red Hat 4.4.6-4)] on linux2
Type "help", "copyright", "credits" or "license" for more information.
(InteractiveConsole)
>>> from __future__ import unicode_literals
>>> import struct
>>> struct.unpack('>I', '\x00\x00\x00\x07')
(7,)
>>>
```

WITHOUT IPYTHON

IS DJANGO MONKEY
PATCHING STRUCT?

No it wasn't...

3 STATEMENTS

1. This **should** break, **every time**.
2. It **breaks** consistently outside of Django env.
3. It **works** consistently within Django env.

The background is a heavily blurred, low-angle shot of a person's face and hand. The person appears to be looking upwards and slightly to the right. The hand is visible on the left side, with fingers slightly curled. The overall tone is soft and out of focus, with a light blue and grey color palette.

SOMEONE DISPROVED 3


```
>>> import struct
>>> struct.unpack('>I', '\x00\x00\x00\x07')
(7,)
>>> struct.unpack(u'>I', '\x00\x00\x00\x07')
(7,)
```

WHAT IS GOING ON?

IT ONLY FAILS IF THE FIRST
CALL PASSES IN UNICODE

```
>>> import struct
>>> struct.unpack(u'>I', '\x00\x00\x00\x07')
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
TypeError: Struct() argument 1 must be string, not unicode
>>> struct.unpack('>I', '\x00\x00\x00\x07')
(7,)
>>> struct.unpack(u'>I', '\x00\x00\x00\x07')
(7,)
```

CONFIRMED

WAS THERE SOME SORT OF
CACHE?

From https://hg.python.org/cpython/file/3a1db0d2747e/Modules/_struct.c


```
1 #define MAXCACHE 100
2 static PyObject *cache = NULL;
3
4 cache_struct(PyObject *fmt)
5 {
6     ...
7     if (s_object != NULL) {
8         if (PyDict_Size(cache) >= MAXCACHE)
9             PyDict_Clear(cache);
10         /* Attempt to cache the result */
11         if (PyDict_SetItem(cache, fmt, s_object) == -1)
12             PyErr_Clear();
13     }
14     return s_object;
15 }
```

YES!

AND THERE'S A BUG REPORT

- From <https://bugs.python.org/issue19099>

“Struct constructor accepts only `str` and not `unicode`. But `struct.pack()` uses caching and it found `Struct('B')` in the cache (because `u'B'` and `'B'` are equal and have same hash).”




```
$ pwd  
/data/virtualenv/vmportal/lib  
$ grep -r "'>I'" * | wc -l  
77
```

'>I' IS USED AT LEAST 77 TIMES

PUTTING IT ALL TOGETHER

- On Gunicorn boot-up, `Struct ('>I')` is called somewhere, and is cached.
- Our code works.
- Then the cache runs out.
- Under the right circumstances, our code will be called next, causing all subsequent calls to fail (and not be cached)...
- Stays broken until this worker processes other code that calls `Struct ('>I')`.
- Then everything works again.

```
1 try:
2     data_b64 = base64.decodestring(key)
3     int_len = 4
4     str_len = struct.unpack(str('>I'), data_b64[:int_len])[0]
5     if str_len == 7:
6         key_type = data_b64[int_len:int_len + str_len]
7 except Exception:
8     message = 'Content is not a valid RSA or DSA SSH Key.'
9     logger.exception(message, extra={'pid': os.getpid()})
10    raise ValidationError(message)
```

SOLUTION WAS SIMPLE

Type-cast the argument to `str`

BUG FIX WAS IMPLEMENTED IN 2.7.7

- The very next version (<https://hg.python.org/cpython/raw-file/f89216059edf/Misc/NEWS>)
 - Issue #19099: The struct module now supports Unicode format strings.

LESSONS LEARNED

- When things don't make sense: Moar Logging.
- Keep questioning your assumptions. Always try to replicate.
- Be wary of:

```
from __future__ import unicode_literals
```


QUESTIONS?

Psst... we're hiring!



VM FARMS

@vmfarms

By Hany Fahim
Founder and CEO
@iHandroid